



SelfGNN with Hard Negative Sampling

Project by Faith Cramer, Jeffrey Yuan, Zihan Zhao, Sharon Lin





Table of contents

01

Introduction

02

Objectives

03

Methodology

04

**Experiment
Setup**

05

Results

06

Conclusion





01

Introduction



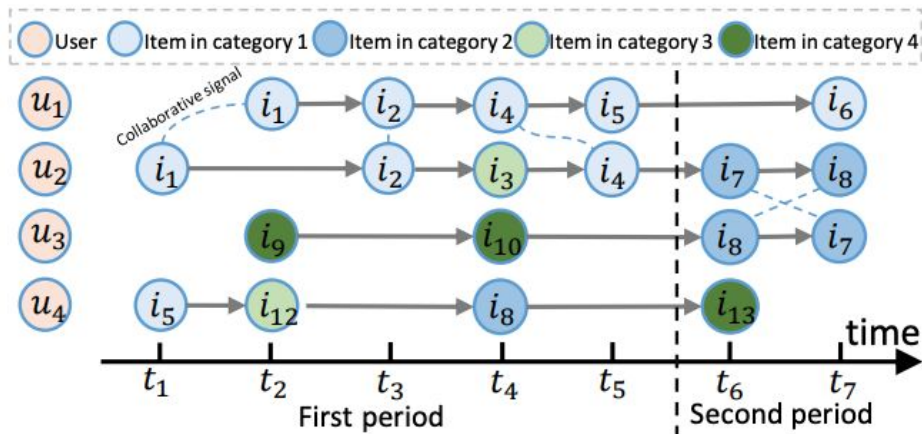


Motivation

- **Sequential recommendation** aims to predict users' future behavior based on their historical interaction patterns.
- Traditional sequential recommenders focus on **single user sequences**, overlooking **collaborative patterns** between users.
- **Conventional** GNNs operate on **static data**, and are incapable of distinguishing between short-term and long-term user behaviors
- Existing self-supervised learning (SSL) methods **depends heavily on high-quality data**, while real-world data is often noisy
- The need to integrate **temporal dynamics**, **collaborative signals**, and **self-supervised learning** under noisy conditions.

What is SelfGNN?

- Novel framework* that captures dynamic user interests
- Combine instance-level **collaborative graph learning** with interval-level **sequential attention modeling**
- Demonstrate consistent and robust **performance gains** over SOTA baselines across multiple datasets

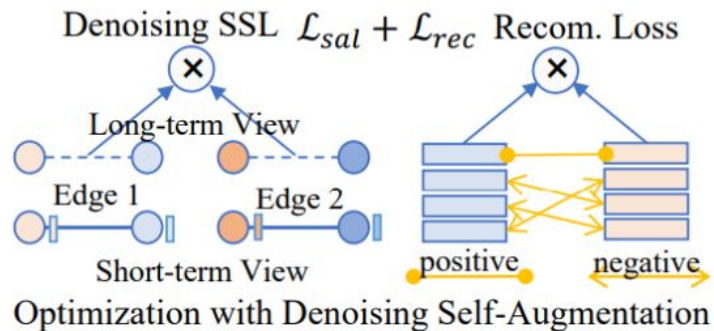

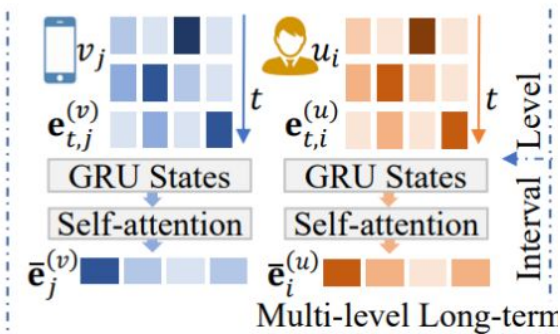
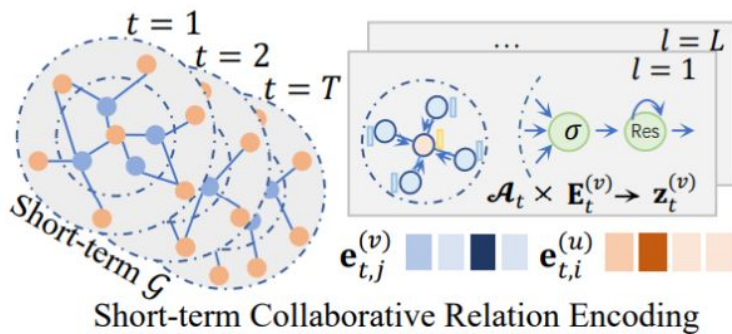


*Liu, Y., Xia, L. and Huang, C., 2024, July. Selfgnn: Self-supervised graph neural networks for sequential recommendation.



Key Components of SelfGNN

- **Short-term Collaborative Graphs Encoding:** constructs time-based user-item graphs and uses GNNs (LightGCN) to capture high-order collaborative patterns in recent user behavior
- **Multi-level Long-term Sequential Learning:** models long-term user preferences at interval and instance levels using GRU and self-attention, enabling temporally-aware user representations
- **Personalized Self-Augmented Learning for Denoising:** Aligns short- and long-term predictions using a user-specific denoising to suppress noisy short-term behaviors and enhance robustness





02

objectives





Improvements to SelfGNN

- Liu et al. highlights that future extensions of SelfGNN should aim to more accurately capture **short-term behavioral characteristics** to improve recommendation performance
- Current framework reduces user-specific noise but does not explicitly distinguish **similar but non-interacted items** from items that users interact with
- Use **contrastive learning**, specifically **hard negative sampling**, to train model to distinguish between relevant / similar but irrelevant items, enhancing precision of short-term user representations



Hard Negative Sampling

- Selects top-K non-interacted items most similar to a user's short-term embedding, using **cosine similarity** to prioritize **informative negatives** over random ones
- Incorporate **InfoNCE contrastive loss** into model's loss function to train the model to pull the user's short-term embedding closer to the true next item (positive), while **pushing it away from highly similar but non-interacted items (hard negatives)**



Project Objectives

- Familiarize ourselves with SelfGNN codebase and **replicate results** of Liu et al. on Amazon, Gowalla, Movielens, and Yelp datasets
- Develop **HardGNN** by integrating **hard negative sampling** into the SelfGNN framework to strengthen the discriminative power of short-term user representations
- **Tune hyperparameters** related to hard negative sampling and **evaluate performance** of HardGNN on all four datasets
- Conduct **ablation studies** to analyze the impact of individual hyperparameter settings on model performance



03

Methodology





Hard Negative Sampling Strategy

- Given user u and positive items i find **negative items j**
- Step 1: Obtain Relevant Embeddings
 - Learned preference embeddings e_u and item embeddings e_j
- Step 2: Calculate Similarity
 - For user preferences, calculate similarity with negative items

$$\text{sim}(e_u, e_j) = \frac{e_u \cdot e_j}{\|e_u\| \|e_j\|}$$

- Step 3: Select Top K Hardest Negatives
 - Rank negatives by similarity to positive items
 - Select Top-K items with highest similarity score



InfoNCE Contrastive Loss

- Contrastive learning is a self-supervised learning technique that aims to pull semantically similar (positive) samples together and push dissimilar (negative) samples apart in the embedding space
- InfoNCE is a popular contrastive loss function that distinguishes positive samples from a set of negative samples

$$\mathcal{L}_{\text{InfoNCE}}(\mathbf{q}, \mathbf{k}^+, \{\mathbf{k}_i^-\}) = -\log \frac{\exp(\text{sim}(\mathbf{q}, \mathbf{k}^+)/\tau)}{\exp(\text{sim}(\mathbf{q}, \mathbf{k}^+)/\tau) + \sum_{j=1}^{N-1} \exp(\text{sim}(\mathbf{q}, \mathbf{k}_j^-)/\tau)}$$

- Anchor (\mathbf{q}): User preference embedding from their sequence (\mathbf{e}_u)
- Positive Key (\mathbf{k}^+): Embedding of the target positive item (\mathbf{e}_i)
- Negative Keys $\{\mathbf{k}_i^-\}$: Embeddings of negative items (\mathbf{e}_j)
- Tau (τ): Temperature Hyperparameter



Integrating HNS and InfoNCE

- The original SelfGNN loss is augmented with the InfoNCE loss

$$\mathbf{L}_{\text{Total}} = \mathbf{L}_{\text{Rec}} + \lambda \cdot \mathbf{L}_{\text{InfoNCE}}$$

- Lambda (λ): Contrastive loss weight parameter

Research Questions

- How does HardGNN perform *w.r.t* top- k recommendation as compared with SelfGNN?
- What are the benefits of the components proposed in HardGNN?
- How do contrastive loss weight and number of hard negative samples impact HardGNN?



04

Experiment Settings





Datasets

- **Amazon-book**: users' ratings of Amazon books in 2014
- **Gowalla**: user geolocation check-in dataset from Gowalla in 2009-10
- **Movielens**: users' ratings for movies from 2002 to 2009
- Unable to process **Yelp** dataset from source code given
- Split each dataset into test set (most recent interaction of each user), validation set (second most recent interaction), and training set (remaining interactions)



Evaluation Metrics

- **Hit Rate (HR)**: measures whether the true next item appears in the top-N predicted results
- **Normalized Discounted Cumulative Gain (NDCG)**: measures how high true item ranks in top-N predicted results, rewarding higher placements more
- Use $N = 10$ and 20



05

Results





Experimental Setup

- Perform **grid search** over number of hard negatives (3, 5, 7) and contrastive loss weights (0.2, 0.1, 0.01)
- Train each configuration for 25 epochs to efficiently **identify best-performing combination** for each dataset
- **Retrain the model** using selected hyperparameters for 100 epochs
- Evaluate **final model performance** on both top-10 and top-20 recommendations and compare with SelfGNN



Results

100 Training Epochs

Dataset	Amazon				Gowalla				Movielens			
Top N	Top 10		Top 20		Top 10		Top 20		Top 10		Top 20	
Metric	HR	NDCG	HR	NDCG	HR	NDCG	HR	NDCG	HR	NDCG	HR	NDCG
<i>SelfGNN</i>	0.371	0.218	0.486	0.250	0.568	0.370	0.703	0.407	0.224	0.121	0.326	0.145
<i>HardGNN</i>	0.223	0.120	0.401	0.203	0.624	0.408	0.736	0.434	0.206	0.111	0.308	0.137

- HardGNN trained with best hyperparameters for each dataset (from gridsearch)
 - Amazon: $K = 3$, $\lambda = 0.1$
 - Gowalla: $K = 3$, $\lambda = 0.2$
 - Movielens: $K = 3$, $\lambda = 0.01$



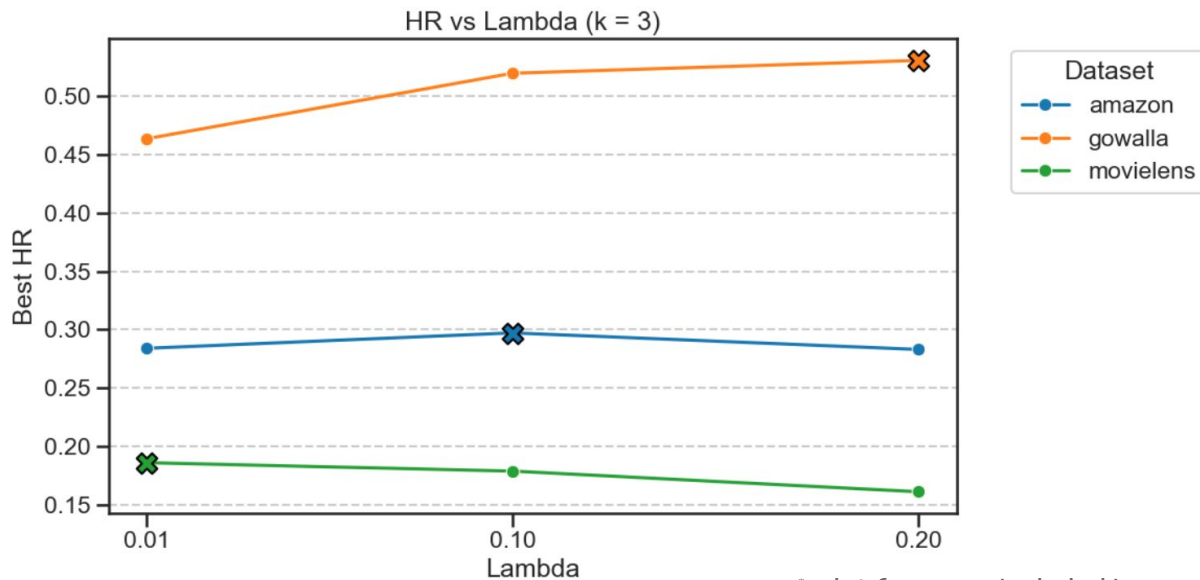
Ablation Studies

- Use the results from the grid search to investigate model performance across **varying values of the two hard negative sampling parameters**
 - The number of hard negatives $k = (3, 5, 7)$ and the contrastive loss weight $\lambda = (0.01, 0.1, 0.2)$
- Examine how performance changes as K varies while holding λ constant
- Examine how performance changes as λ values while holding K constant



Effect of Contrastive Loss Weight

25 Training Epochs

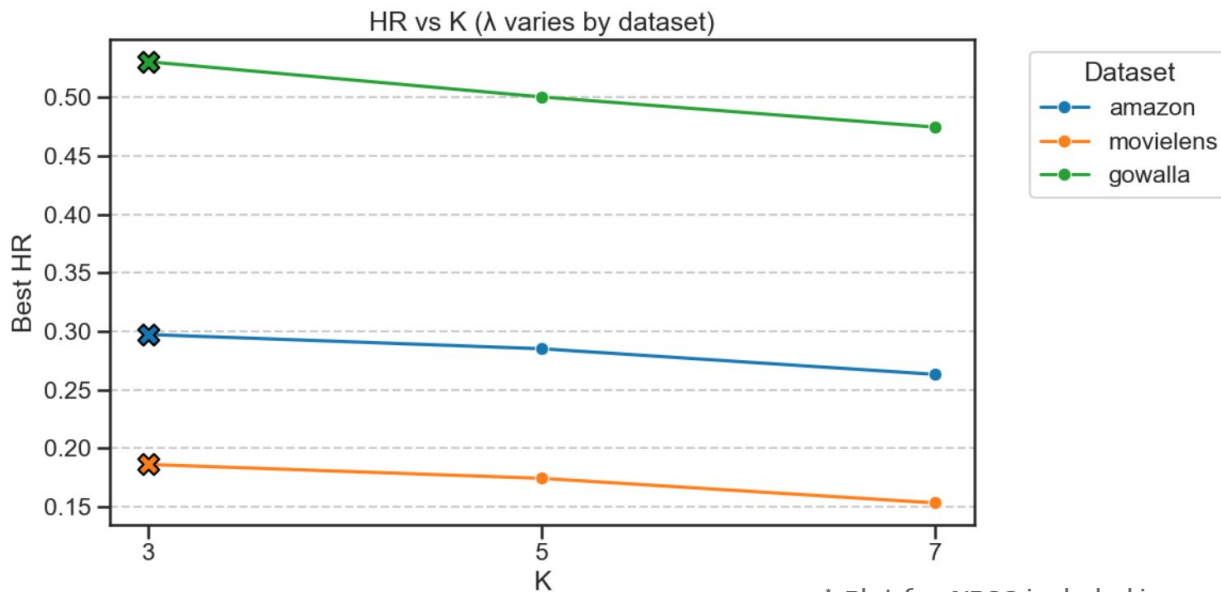


* Plot for NDCG included in appendix



Effect of # Hard Negative Samples

25 Training Epochs



* Plot for NDCG included in appendix



06

conclusion





Summary of Findings

- HardGNN (our solution) **underperformed** SelfGNN (original work) for Amazon and Movielens datasets
 - These datasets rely more on long-term signals, so emphasizing short-term behavior via hard negatives may disrupt model balance and hurt performance
- HardGNN **outperformed** Self GNN on Gowalla dataset, likely reflecting the benefits of hard negative sampling for the specific characteristics of this dataset
 - Gowalla is the sparsest dataset, where random negatives (e.g., distant or irrelevant locations) rarely offer meaningful learning signals
 - Hard negatives, being contextually similar but behaviorally distinct, provide more informative information for training



Future Work

- Improve the time and space efficiency of HardGNN implementation
- Conduct more extensive hyperparameter tuning and run grid search for more training epochs
- Train HardGNN to 150 epochs (instead of 100) and retrain 5 times (instead of once) to obtain p-values to align with original SelfGNN experimental setup



Thanks!
Questions?





Appendix

